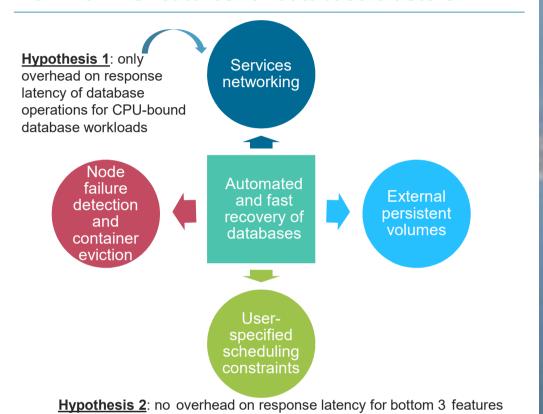


Evaluation of container orchestration systems for DistriN=t deploying and managing NoSQL database clusters

Container orchestration (CO) systems, such as Docker Swarm, Kubernetes, Mesos Marathon and DC/OS, have been initially designed for stateless services. However, they have also been used for running database clusters due to improved automated management. This work evaluates the performance overhead of Docker Swarm and Kubernetes when running NoSQL database clusters, with MongoDB and Cassandra as case study.

https://github.com/eddytruyen/containers_on_openstack/blob/master/README.md IEEE International Conference on Cloud Computing (IEEE Cloud 2018). Workshop 2: Cloud Infrastructure, Tuesday July 3, 8:30am

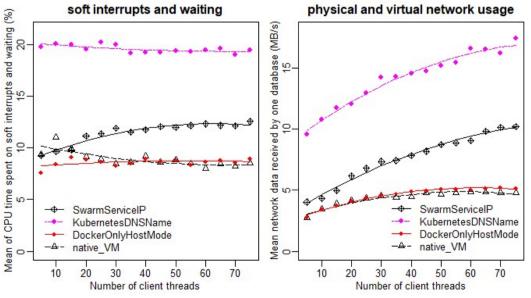
Common CO features for database clusters



Results for YCSB benchmark[1]

Hypothesis 1:

For CPU-bound Cassandra workloads: services networking causes overhead due to intermediate virtual network bridge in comparison to DockerOnly without NAT (--net=host option)



Resource usage graphs for Cassandra 2.0; 3-node cluster, Openstack VMs 2CPU 4GB, Ubuntu 16.04; YCSB workload A, 100% insert operations, size of dataset = 11.7GB => write-optimized database

For CPU-bound MongoDB workloads: Wrt the mean and 95th quantile of response latency metrics, common rankings between services networking approaches are found for specific combinations of YCSB workload type and operation type:

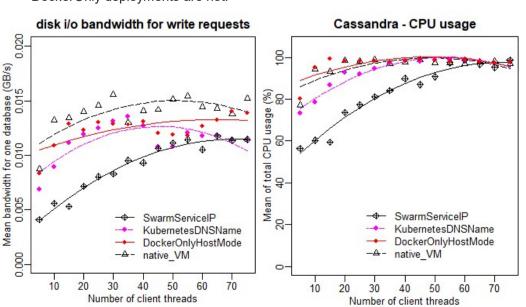
Three services networking approaches

Service networking approach	Common features among CO systems	Swarm	Kubernetes	Marathon	DC/OS
Routing mesh with service node ports	Distributed Layer 4 Load balancer (with ipvs)	✓	✓		
	Central L4-L7 load balancer (without ipvs)	✓	✓	✓	✓
Virtual network with service IPs	Distributed Layer4 load balancer (with ipvs)	✓	✓		\checkmark
	with stable DNS name for services	✓	✓		✓
	IP per container	✓	✓	✓	✓
Host ports networking with VM IPs	Mapping container port to host port	✓	✓	✓	√
	Automated allocation of host ports	✓		✓	\checkmark
	Static host port conflict management	√	✓		

Operation	Update-heavy workloads (A, F)	Read-heavy workloads (B,C,D)
Read	4.8,SwarmFloatingIPHostPort 4.9,SwarmHostNameNodePort 5.5,K8FloatingIPHostPort 5.5,K8ServiceIP 5.6,K8HostNameNodePort 5.7,K8ContainerDNSName 23.6,SwarmServiceIP	6.7,K8ServiceIP 7.1,K8ContainerDNSName 7.5,SwarmFloatingIPHostPort 7.9,K8HostNameNodePort 7.9,SwarmHostNameNodePort 8.2,K8FloatingIPHostPort 10.6,SwarmServiceIP
Update/ Insert	20.2,SwarmFloatingIPHostPort 21.1,K8FloatingIPHostPort 21.3,K8ServiceIP 21.8,K8HostNameNodePort 22.2,K8ContainerDNSName 23.8,SwarmHostNameNodePort 43.8,SwarmServiceIP	18.9,K8HostNameNodePort 19.1,K8FloatingIPHostPort 20.2,SwarmFloatingIPHostPort 21.3,SwarmHostNameNodePort 26.1,K8ContainerDNSName 27,5,K8ServiceIP 37.5,SwarmServiceIP

95th quantile of response latencies for MongoDB 3.2.4; 3 node cluster, VMs 2CPU 4GB,Ubuntu 16.04; size of dataset = 23MB; => <u>in-memory database</u>

Hypothesis 2: For CPU-bound Cassandra workloads: in comparison to VMbased deployments, local volume plugins for Docker Swarm and Kubernetes are disk i/o performance bottlenecks. However, default volume drivers in DockerOnly deployments are not!



Funded by: IWT-DeCoMAdS

[1] B. F. Cooper, A. Silberstein, E. Tam, R. Ramakrishnan, and R. Sears, "Benchmarking cloud serving systems with YCSB," Proc. 1st ACM Symp. Cloud Comput. - SoCC '10, pp. 143-154, 2010.







In collaboration with: Red Hat