

Container orchestration (CO) systems, such as Docker Swarm, Kubernetes, Mesos Marathon and DC/OS, have been initially designed for stateless services. However, they have also been used for running database clusters due to improved automated management. This work evaluates the performance overhead of Docker Swarm and Kubernetes when running NoSQL database clusters, with MongoDB and Cassandra as case study.

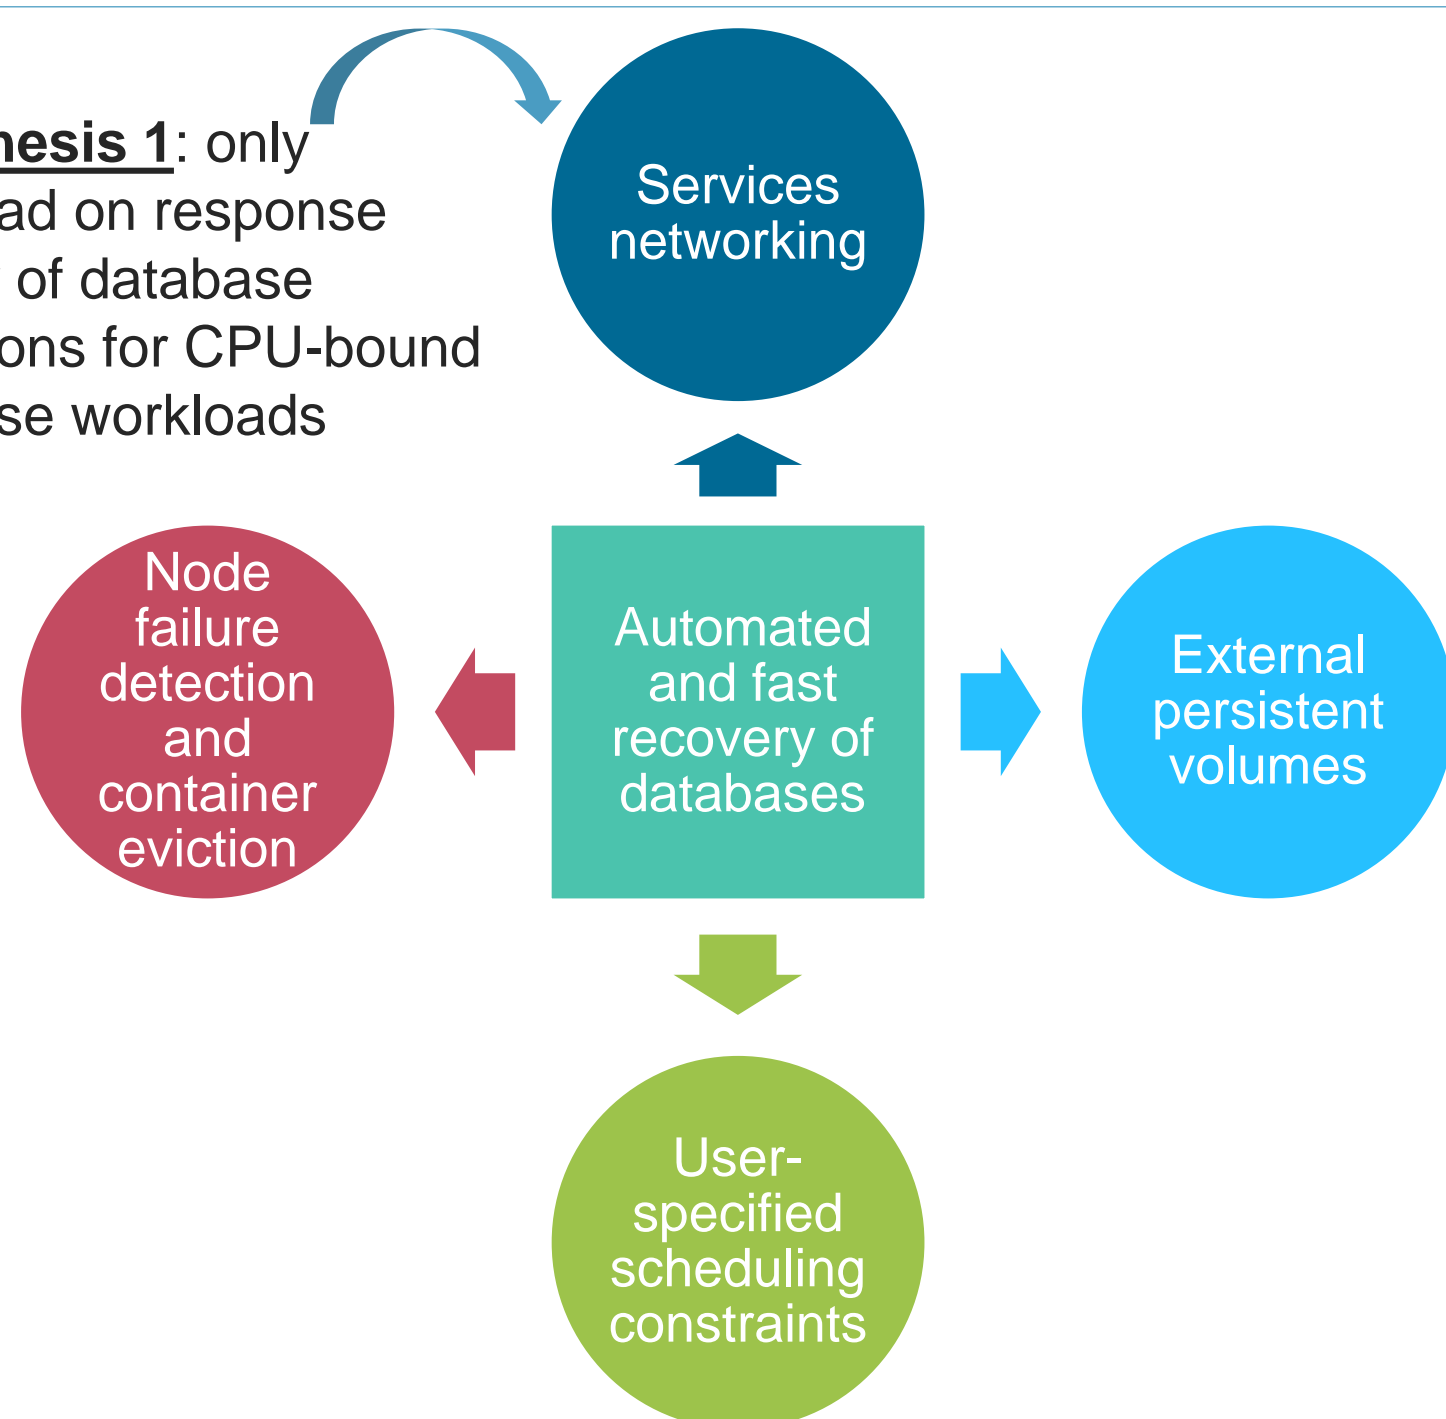
[https://github.com/eddytruyen/containers\\_on\\_openstack/blob/master/README.md](https://github.com/eddytruyen/containers_on_openstack/blob/master/README.md)

IEEE International Conference on Cloud Computing (IEEE Cloud 2018),

Workshop 2: Cloud Infrastructure, Tuesday July 3, 8:30am

## Common CO features for database clusters

**Hypothesis 1:** only overhead on response latency of database operations for CPU-bound database workloads

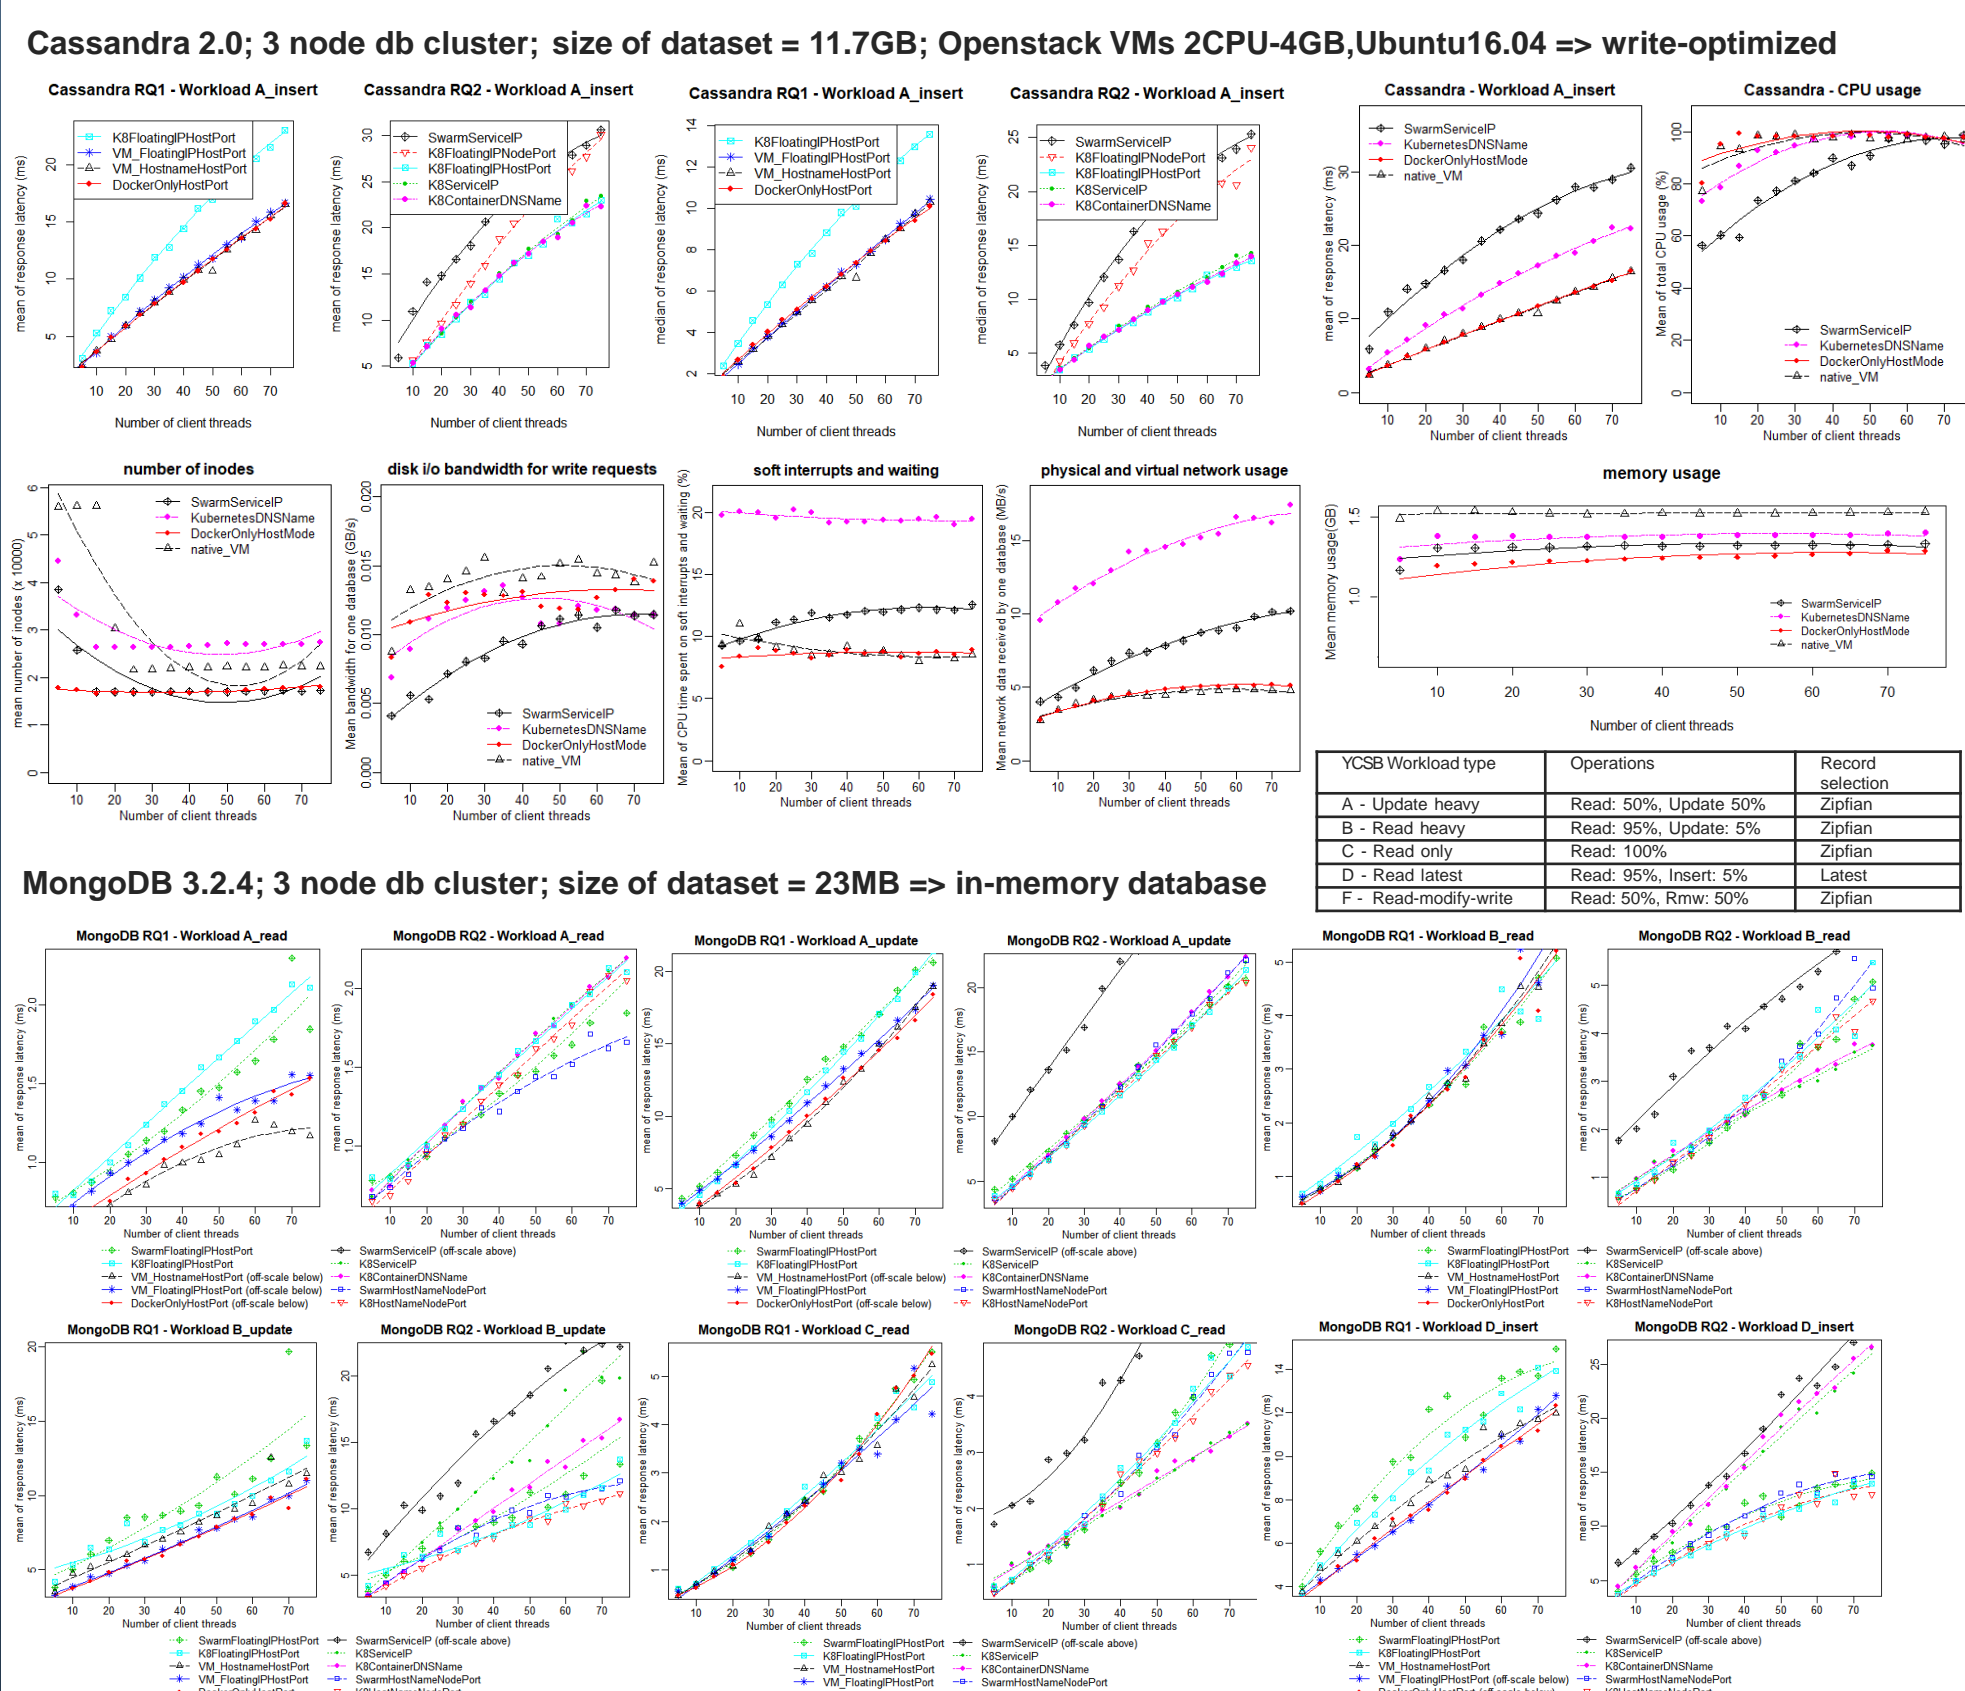


**Hypothesis 2:** no overhead on response latency for bottom 3 features

## Three services networking approaches

Service networking approach	Common features among CO systems	Docker Swarm	Kubernetes	Marathon	DC/OS
Routing mesh for global service ports	Layer4 distributed load balancer (ipvs support)	✓	✓		
	central L4-L7 load balancer (without ipvs)	✓	✓	✓	✓
Virtual IP network for containers	Layer4 distributed load balancer (with ipvs)	✓	✓		✓
	with stable DNS name for services	✓	✓		✓
	IP per container	✓	✓	✓	✓
Host ports networking	Mapping container port to host port	✓	✓	✓	✓
	Automated allocation of host ports	✓		✓	✓
	Static host port conflict management	✓	✓		

## Results for YCSB benchmark<sup>[1]</sup>



## Findings for CPU-bound DB workloads

**Hypothesis 1:**

- Services networking causes overhead due to intermediate virtual network bridge in comparison to DockerOnly without NAT (--net=host option)
  - Run containers of services in host mode (beta feature Swarm)
  - Leverage container security policies to improve security isolation (stable feature K8)
- Common rankings between services networking approaches with respect to 95<sup>th</sup> quantile of response latency for specific combinations of YCSB workload types and operation type

Operation	Update-heavy workloads (A, F)	Read-heavy workloads (B,C,D)
Read	4.5.SwarmFloatingIPHostPort	6.7.K8ServiceIP
	4.9.SwarmHostNameNodePort	7.1.K8ContainerDNSName
	5.3.K8FloatingIPHostPort	7.5.SwarmFloatingIPHostPort
	5.3.K8ServiceIP	7.9.K8HostNameNodePort
	5.6.K8HostNameNodePort	7.9.SwarmHostNameNodePort
	5.7.K8ContainerDNSName	8.2.K8FloatingIPHostPort
Update/Insert	20.2.SwarmFloatingIPHostPort	18.9.K8HostNameNodePort
	21.1.K8FloatingIPHostPort	19.1.K8FloatingIPHostPort
	21.3.K8ServiceIP	20.2.SwarmFloatingIPHostPort
	21.8.K8HostNameNodePort	21.3.SwarmHostNameNodePort
	22.2.K8ContainerDNSName	26.1.K8ContainerDNSName
	23.8.SwarmHostNameNodePort	27.5.K8ServiceIP

each MongoDB deployment is pre-annotated with the mean of the observed 95<sup>th</sup> quantiles of response latency (ms).

**Hypothesis 2:**

- Local volume drivers are disk i/o performance bottlenecks in Docker Swarm and in Kubernetes, but not in DockerOnly deployments
  - As a result, in opposition to DockerOnly and VM-based deployments, 100% CPU utilization cannot be achieved